

# Automated testing of data integration solutions

Slawomir Chodnicki Bl-consultant

slawomir.chodnicki@twineworks.com

# Day 1 overview



- hands-on project intro
- various kinds of automated tests
- testable solutions vs. untestable messes
- minimizing external dependencies
- fixtures & helpers
- declarative tests

# Day 1 overview (continued) Twineworks



- Using ETL to test ETL solutions
  - PDI techniques for unit tests
  - PDI techniques for integration tests
  - practical limitations of ETL-based tests

# Day 2 overview



- Scripting tests
  - jruby as a scripting language
  - scripting helpers for command execution and fixture loading

# Day 2 overview (continued) Twineworks



- Organizing the test suite with rspec
  - rspec output
  - rspec features
  - rspec reports

# Day 3 overview



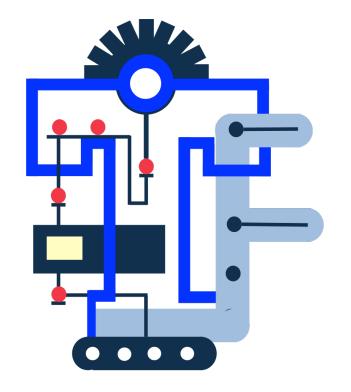
- Continuous Testing with Jenkins
  - installing jenkins
  - access control in jenkins
  - useful jenkins plugins
  - setting up jenkins to run ETL tests
  - test reports

# Day 3 overview (continued) Twineworks



- Deployment
  - deploying from version control
  - migration scripts





Demo project

#### Hands-on project intro



```
|-- bin  # entry point scripts and dependencies
|-- data  # source data ingested by the ETL
|-- environments  # environment configuration
|-- etl  # ETL solution
|-- spec  # ETL test suite
|-- logs  # logging location
|-- spec  # jruby tests and helpers
```

# database configuration



- database: 'dwh'
- user: 'etl'
- password: 'password'

# database configuration



```
$ mysql -u root

CREATE DATABASE dwh;

GRANT ALL PRIVILEGES ON dwh.* TO etl@'%' identified by 'password';

GRANT ALL PRIVILEGES ON dwh.* TO etl@'localhost' identified by 'password';

FLUSH PRIVILEGES;
```

# project configuration



- environment configuration in
  - environments/local
    - environment
    - my.cnf
    - .kettle/shared.xml
    - .kettle/kettle.properties

#### environment



```
#!/bin/bash
export ROBOT_DB_NAME=dwh
export ROBOT_LOG_BASE_DIR=/Users/slawo/Desktop/etl-testing/logs
export ROBOT_PDI_HOME=/Users/slawo/pentaho/data-integration-5.4
```

# my.cnf



```
[client]
user=etl
password=password
host=localhost
```

# .kettle/kettle.properties



```
### Settings for Robot
ENV_MACHINE=local
ENV_DATA_DIR=/Users/slawo/Desktop/etl-testing/solution/data/in
```

# setting hostname explicitly shortens PDI startup time
KETTLE\_SYSTEM\_HOSTNAME=localhost

#### .kettle/shared.xml



```
<?xml version="1.0" encoding="UTF-8"?>
<sharedobjects>
    <connection>
         <name>dwh</name>
         <server>localhost</server>
         <type>MYSQL</type>
         <access>Native</access>
         <database>dwh</database>
         <port>3306</port>
         <username>etl</username>
         <password>Encrypted 2be98afc86aa7f2e4bb18bd63c99dbdde/password>
    </connection>
</sharedobjects>
```

#### Initialize db



```
$ bin/robot db reset
clearing database [ OK ]
initializing database
2017/06/20 15:32:37 - Kitchen - Start of run.
2017/06/20 15:32:37 - reset_dwh - Start of job execution
2017/06/20 15:32:41 - Kitchen - Finished!
2017/06/20 15:32:41 - Kitchen - Processing ended after 4 seconds.
\Gamma OK \Gamma
```

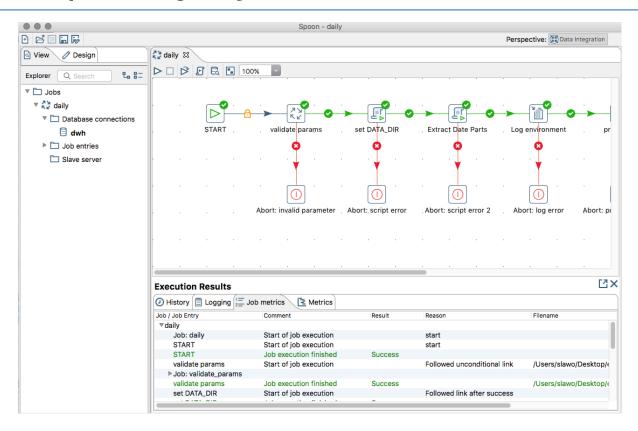
# Run daily



\$ bin/robot spoon

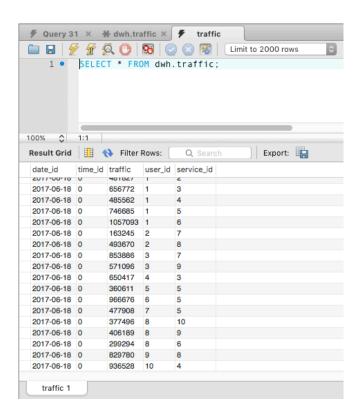
# Run etl/daily.kjb



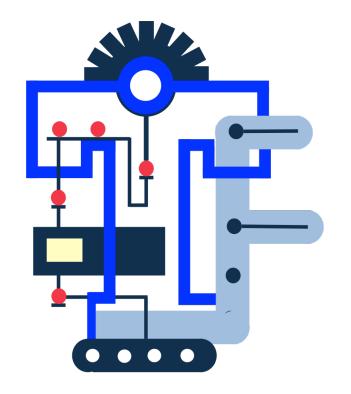


#### Inspect database contents









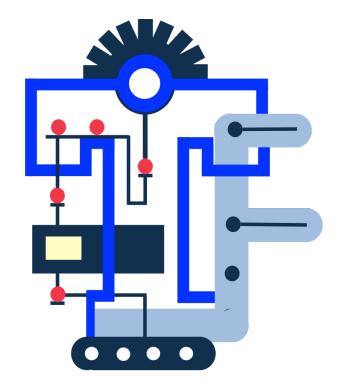
Testing ETL solutions

#### Kinds of automated tests



- Unit tests
- Integration tests
- Functional tests
- Non-functional tests





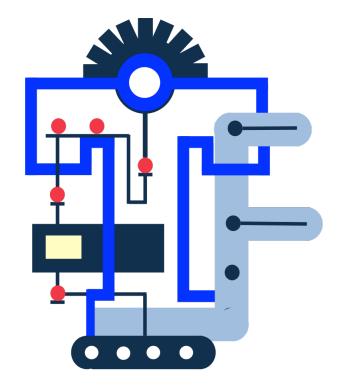
Unit tests

#### Unit tests



- Single isolated unit of ETL under test
- A unit performs a computation (no side-effects)
- What is a "unit" in PDI?
  - Job?
  - Transformation?
  - Sub-transformation (mapping)?





Testing jobs

#### A simple unit test



Create a job
etl/spec/dwh/validate\_params/validate\_params\_spec.kjb

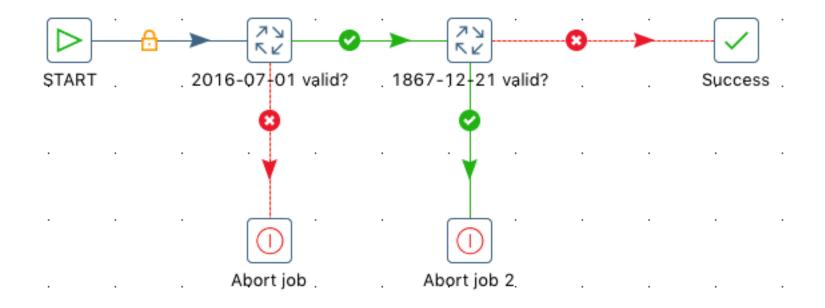
```
The job calls etl/dwh/validate_params.kjb
```

- with DATA\_DATE=2016-07-01 and expects it to succeed
- with DATA\_DATE=1867-12-21 and expects it to fail

The job succeeds if all expectations are met. It fails otherwise.

#### A simple unit test





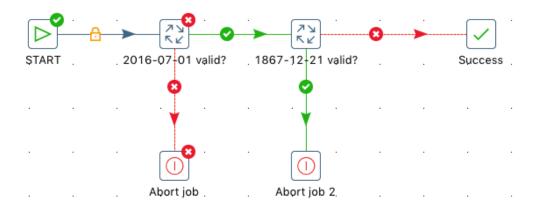
#### Test must run in known env



\$ bin/robot db clear

Run

etl/spec/dwh/validate\_params/validate\_params\_spec.kjb



#### Environment reset



Create a job
etl/spec/support/reset\_all.kjb

The job runs

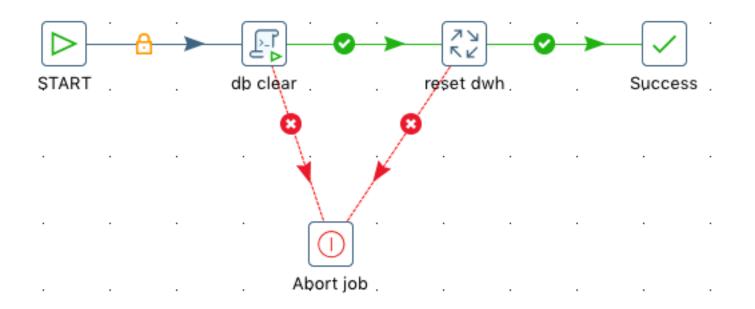
\$ bin/robot db clear

Then calls
etl/init/dwh/reset\_dwh.kjb

The job succeeds if both tasks succeed. It fails otherwise.

#### Environment reset

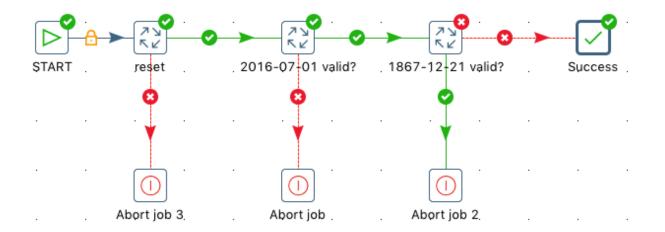




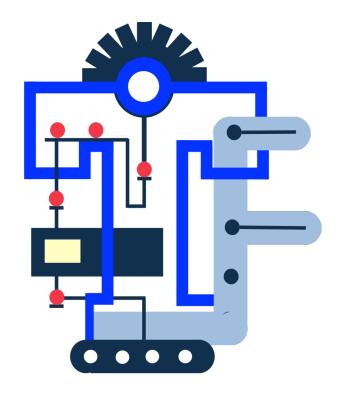
#### Test must run in known env



Fix
etl/spec/dwh/validate\_params/validate\_params\_spec.kjb
to establish a known environment before testing.







Testing transformation results

#### Test transformation results Twineworks



Have a look at etl/util/day\_sequence.ktr

Rows of step: read dates (5 rows)

| # ^ | date_id    | year | month | day |
|-----|------------|------|-------|-----|
| 1   | 2000-01-01 | 2000 | 1     | 1   |
| 2   | 2000-01-02 | 2000 | 1     | 2   |
| 3   | 2000-01-03 | 2000 | 1     | 3   |
| 4   | 2000-01-04 | 2000 | 1     | 4   |
| 5   | 2000-01-05 | 2000 | 1     | 5   |
|     |            |      |       |     |
|     |            |      |       |     |
|     |            |      |       |     |

#### Test transformation results



```
Create a job
etl/spec/util/day_sequence/day_sequence_spec.kjb
```

```
The job calls

etl/util/day_sequence.ktr

With parameters

START_DATE=2017-02-27

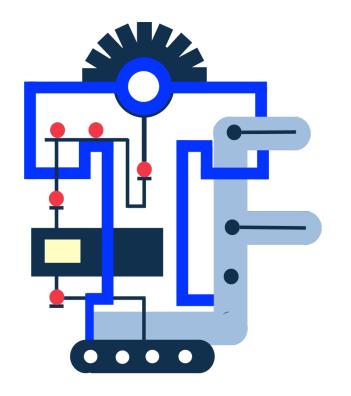
DAY_COUNT=3
```

The job succeeds if and only if the transformation returns:

- exactly 3 rows, representing 2017-02-27, 2017-02-28, 2017-03-01
- with at least 4 fields each
  - string date\_id
  - int year
  - int month
  - int day

Hint: the JavaScript job entry exposes result rows as 'rows'. http://wiki.pentaho.com/display/EAI/JavaScript+%28job+entry%29





Testing sub-transformations

#### Test sub-transformations



Have a look at

etl/util/string\_cleaner.ktr

It transforms a string to be a valid identifier in an internal system. The system accepts identifiers of the following form:

Only ascii letters a-z, A-Z, digits 0-9 and the underscore \_ are allowed. The identifier may not start with a digit.

#### Test sub-transformations



Conversion is specified as follows:

Diacritics are reduced to the base character. So Ä becomes A, é becomes e, etc. Strings beginning with a digit are prefixed with an underscore character.

The following characters are replaced by letter sequences:

```
& _and_ $ _dollar_
* _star_ £ _pound_
- _dash_ € _euro_
% _percent_ . _dot_
@ _at_ , _comma_
# hashtag
```

Any other invalid character is replaced with an underscore character.



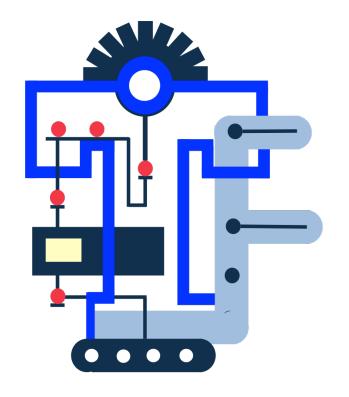
Create a job

/etl/spec/util/string\_cleaner/string\_cleaner\_spec.kjb

The job verifies that **etl/util/string\_cleaner.ktr** performs the conversion as per the input and expected columns of /etl/spec/util/string\_cleaner/string\_cleaner\_spec.csv

The csv file is UTF-8 encoded.





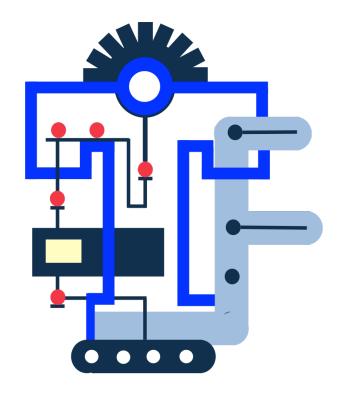
Integration tests

# Integration tests



- ETL responsible for a set of related sideeffects under test
- Most common case in ETL testing
  - Test individual phases of a batch process





Tests for the pre-flight phase

# Test the pre-flight phase



Have a look at

/etl/dwh/pre\_flight/pre\_flight.kjb

Given a DATA\_DIR, DATA\_DATE, as well as the corresponding DATA\_YEAR, DATA\_MONTH, and DATA\_DAY. It verifies that a non-empty file exists at location:

DATA\_DIR/DATA\_YEAR/DATA\_MONTH/DATA\_YEAR-DATA\_MONTH-DATA\_DAY.csv



Create a job

etl/spec/dwh/pre\_flight/pre\_flight\_ok\_spec.kjb

The job verifies that **/etl/dwh/pre\_flight/pre\_flight.kjb** succeeds when the file specified by the paramaters exists.

Use the folder /spec/fixtures/pre\_flight to act as DATA\_DIR to host any test files.



Create a job

etl/spec/dwh/pre\_flight/pre\_flight\_empty\_file\_spec.kjb

The job verifies that **/etl/dwh/pre\_flight/pre\_flight.kjb** fails when the file specified by the paramaters exists, but is empty.

Use the folder /spec/fixtures/pre\_flight to act as DATA\_DIR to host any test files.

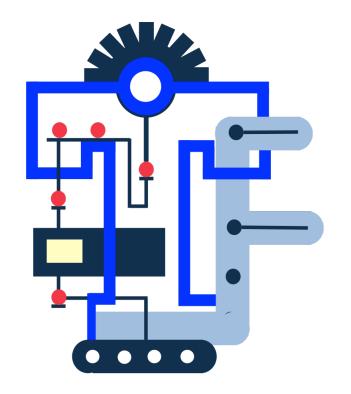


Create a job

etl/spec/dwh/pre\_flight/pre\_flight\_missing\_file\_spec.kjb

The job verifies that **/etl/dwh/pre\_flight/pre\_flight.kjb** fails when the file specified by the paramaters does not exist.





Tests for the stage phase

# Test the stage phase



Have a look at

/etl/dwh/stage/stage.kjb

Given a DATA\_DIR, DATA\_DATE, as well as the corresponding DATA\_YEAR, DATA\_MONTH, and DATA\_DAY. It loads the contents of the file DATA\_DIR/DATA\_YEAR/DATA\_MONTH/DATA\_YEAR-DATA\_MONTH-DATA\_DAY.csv into the database table: stage\_traffic.

## Test the stage phase



Create a job

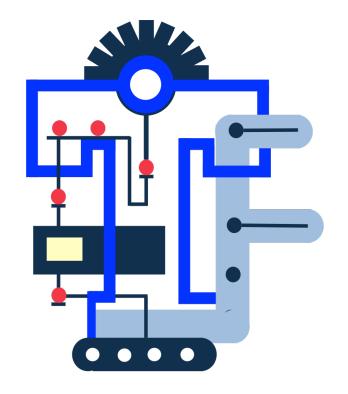
etl/spec/dwh/stage/stage\_spec.kjb

The job verifies that /etl/dwh/stage/stage.kjb

- succeeds when a valid file is loaded
- table **stage\_traffic** contains the expected number of records
- table **stage\_traffic** contains one specific record from the loaded file

Use the folder /spec/fixtures/stage to act as DATA\_DIR to host any test files.







#### Have a look at

#### /fixtures/load/2017-06-21.json

```
"stage_traffic": [
    "date_id": "2017-06-21",
    "time id": 4,
    "user": "Gulliver",
    "service": "Libslack",
    "traffic": 100
    "date_id": "2017-06-21",
    "time_id": 4,
    "user": "Gulliver",
    "service": "Libslack",
    "traffic": 200
 },
```



```
"table_name_1": [
    "field 1": "value 1",
    "field 2": value 2,
    "field 3": "value 3",
    "field_4": "value_4",
    "field 5": value 5
  },
    "field 1": "value 1",
    "field_2": value_2,
    "field 3": "value 3",
    "field 4": "value 4",
    "field 5": value 5
  },
```

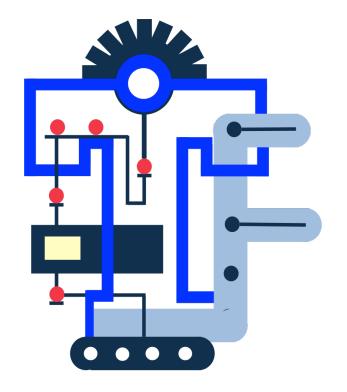


Create a transformation

etl/spec/support/load\_json\_fixture.ktr

Given a JSON file of the above structure in parameter INPUT\_FILE, it loads the data from the json file into the dwh database.





Test the load phase

## Test the load phase



Have a look at

/etl/dwh/load/load.kjb

Given a DATA\_DATE, it deletes any records with that date\_id from the traffic table, then aggregates records from **stage\_traffic**, calculates the user\_tag field, populates dimension tables **user**, and **service**, and writes fact records to fact table **traffic**.

#### Test the load phase



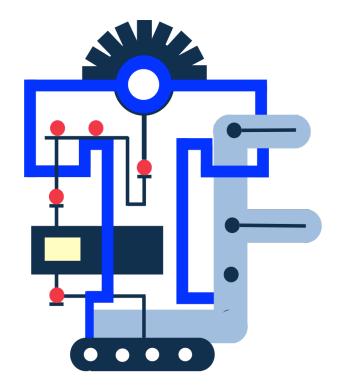
Create a job etl/spec/dwh/load/load\_spec.kjb

The job verifies that /etl/dwh/load/load.kjb

- loads data from stage\_traffic
- has aggregated the records for same date, time, user and service
- has calculated the user\_tag by transforming it as expected
- has filled the service table as expected

Use /fixtures/load/2017-06-21.json to populate the stage\_traffic table.





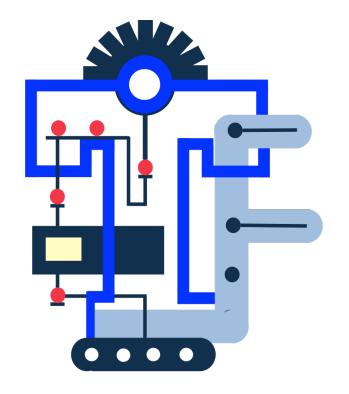
Functional testing

#### Functional tests



- Entry point of ETL solution under test
- Assertions reflect functional contract
  - Behavior on happy path
  - Behavior on errors
  - Behavior on incorrect invocation





A comparison helper

## A comparison helper



Have a look at

/spec/fixtures/functional/expected/2017-06-20.csv

It contains the following columns

| Name     | Туре    |
|----------|---------|
| date_id  | String  |
| time_id  | Integer |
| user     | String  |
| user_tag | String  |
| service  | String  |
| traffic  | Integer |

This structure corresponds a full join of all data warehouse traffic information. Dimension keys are omitted.

#### A comparison helper



Create a transformation

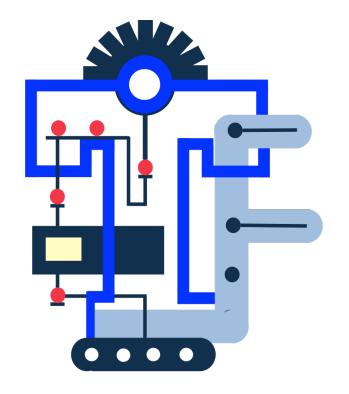
etl/spec/support/compare\_state\_traffic.ktr

Given a csv file of the above structure in parameter EXPECTED\_FILE, it compares the data from the csv file with the contents of the dwh database.

The transformation succeeds when the contents are equal.

The transformation fails when the contents are different.







Have a look at

/etl/daily.kjb

Given a DATA\_DATE as well as a DATA\_DIR, it loads the contents of the corresponding file into the data warehouse.

DATA\_DATE must exist in the date dimension table, and the corresponding file must exist in DATA\_DIR.



Create a job etl/spec/daily/daily\_0000\_00\_00\_spec.kjb

The job verifies that **/etl/daily.kjb** fails and does not load any data when invoked with invalid DATA\_DATE 0000-00-00.



Create a job etl/spec/daily/daily\_2001\_01\_01\_spec.kjb

The job verifies that **/etl/daily.kjb** fails and does not load any data when invoked with DATA\_DATE 2001-01-01, for which the data file does not exist.



Create a job etl/spec/daily/daily\_2017\_06\_20\_spec.kjb

The job verifies that **/etl/daily.kjb** succeeds and loads expected data when invoked with DATA\_DATE 2017-06-20.

Use the folder /spec/fixtures/functional as DATA\_DIR and /spec/fixtures/functional/expected/2017-06-20.csv to compare database state.



Create a job

etl/spec/daily/daily\_2017\_06\_20\_reload\_spec.kjb

The job verifies that **/etl/daily.kjb** succeeds and loads expected data when invoked with DATA\_DATE 2017-06-20 twice in a row. The expected data is the same as if the date had been loaded only once.

Use the folder /spec/fixtures/functional as DATA\_DIR and /spec/fixtures/functional/expected/2017-06-20.csv to compare database state.



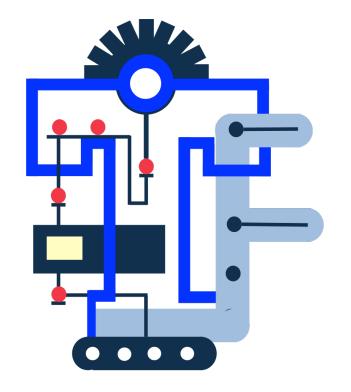
Create a job

etl/spec/daily/daily\_2017\_06\_20\_to\_21\_spec.kjb

The job verifies that **/etl/daily.kjb** succeeds and loads expected data when invoked with DATA\_DATE 2017-06-20, and again with DATA\_DATE 2017-06-21.

Use the folder /spec/fixtures/functional as DATA\_DIR and /spec/fixtures/functional/expected/2017-06-20-to-21.csv to compare database state.





Non-functional tests

#### Non-functional tests



- Performance
  - how long does workload x take?
- Stability
  - what does it take to break it?
  - How much memory is too little?
  - What happens when loading unexpected data? (truncated file, column too long, 50MB XML in string field, badly formatted CSV reads as single field, empty files)

#### Non-functional tests



- Security
  - Verify configuration assumptions automatically
- Compliance
  - We must use version x of library y
  - Often done by inspection

# Test configuration



Create a job

etl/spec/environment/robot\_dir\_is\_set\_spec.kjb

The job verifies that the kettle variable ROBOT\_DIR is set correctly. The variable must point to an existing directory, and below that directory there must be a **bin/robot** file.

The job succeeds if above tests pass. It fails otherwise.

## Test compliance



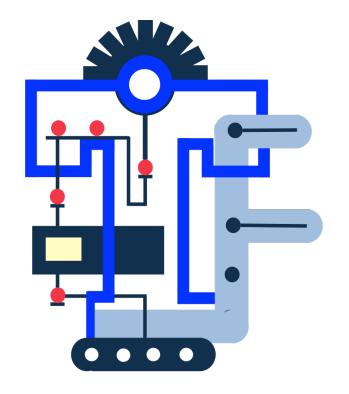
Create a job

etl/spec/environment/mysql\_driver\_spec.kjb

The job verifies that the file **\${ROBOT\_PDI\_HOME}/lib/mysql-connector-java-5.1.42-bin.jar** exists.

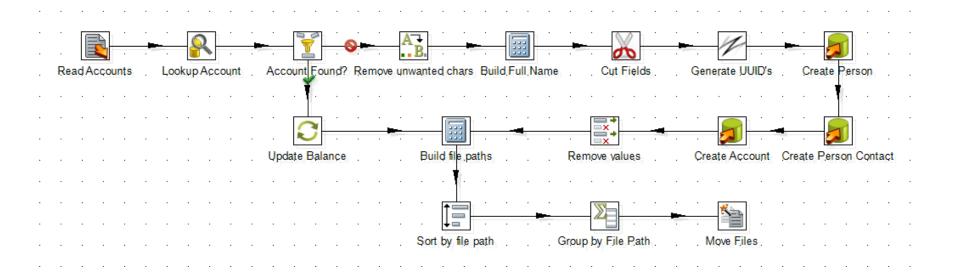
The job succeeds if above test passes. It fails otherwise.



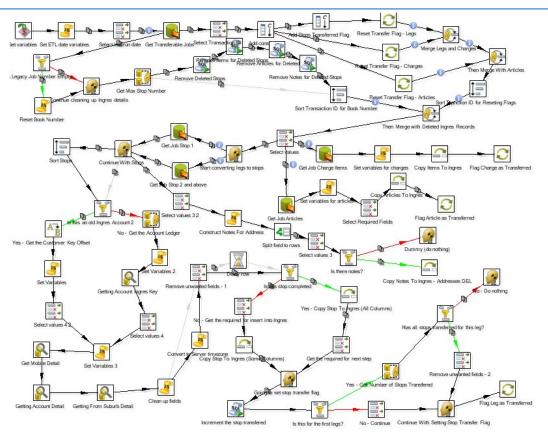


Testing philosophy











- Configuration management
  - configure all data sources/targets and paths through kettle variables or parameters
  - local environment (not in version control)
  - test environment (reference environment)
  - QA environment
  - production environment
- Resist the temptation to share configuration across environments
- Resist the temptation to introduce a hierarchy of environments



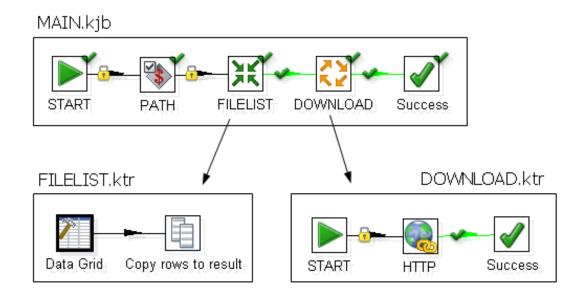
- Define sub-systems/phases
  - define pre-requisites
    - data expected in certain sources
  - define outcomes
    - data written to certain sinks
- A sub-system/phase of the ETL process is responsible for a small set of related side-effects to happen



- Define entry points with a full functional contract.
- An entry point implements an application feature.



Verify behavior of the entity you run directly



### Minimize external deps



- run all data sources and targets locally or through dedicated instances of networked resources
- do not share anything you modify as part of the tests
- if team members cannot run the test suite locally, the test-suite will turn into a chore and a liability
- reproducible results: you can do it, your team mates can do it, and jenkins can do it

## fixtures & helpers



- Data Fixtures
  - sets of test data, encoded in a convenient way, easily loaded into data sources and sinks
- JSON, CSV, SQL, XML, YAML
  - Use whatever is easiest to maintain for the team
- Generate data fixtures through parameterized scripts if you need to generate datasets with consistent relationships

## fixtures & helpers



- File Fixtures
  - sets of test files acted upon during a run
- Maintain file fixtures separate from source location expected by ETL
- If fixture files are changed as part of the test, copy them to a temporary location before running tests
- Create a unique source location per test run, if the file location is shared (like sftp)

## fixtures & helpers



#### Helpers

- utility code/etl of components reused to make tests about the what, not about the how
- fixture loaders
- assertion helpers
- data comparison helpers

#### Declarative tests



 A set of conventions, helpers, fixtures, and generic test runners that make it possible to define a test with minimal configuration and code.

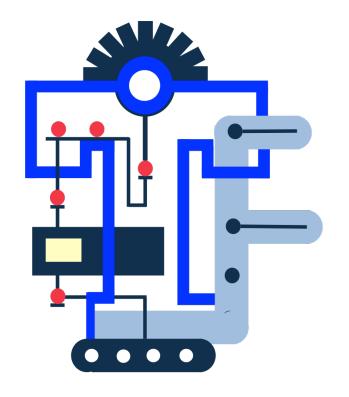
#### **Practical limitations**



- Limited ability to provide test code that is
  - generic
  - data-driven

 metadata injection helps, but the increase in complexity is not trivial





Scripting tests

#### Test orchestration



\$ bin/robot test

```
λ:solution $ bin/robot test
Run options: exclude {:long_running=>true, :remote=>true, :integration=>true}
db clear command
 when db is not empty
   exits with exit code 0
   clears the db
db reset command
 when db is not empty
   exits with exit code 0
   clears the db
   creates a dim date dimension
     including 2000-01-01
     including 2015-12-31
   creates a dim_time dimension
     including 00 AM
     including 11 PM
     including -1/NA
   creates an iso_countries dimension
     including US
     including DE
ETL
 etl/spec/daily/daily_0000_00_00_spec.kjb
   completes successfully
 etl/spec/daily/daily_2001_01_01_spec.kjb
   completes successfully
 etl/spec/daily/daily_2017_06_20_reload_spec.kjb
   completes successfully
 etl/spec/daily/daily_2017_06_20_spec.kjb
```

```
completes successfully
etl/spec/daily/daily_2017_06_20_to_21_spec.kjb
  completes successfully
etl/spec/dummy/dummy_spec.kjb - verifies the test suite runs ETL as tests
 completes successfully
etl/spec/dwh/load/load_spec.kjb
 completes successfully
etl/spec/dwh/pre_flight/pre_flight_empty_file_spec.kjb
  completes successfully
etl/spec/dwh/pre_flight/pre_flight_missing_file_spec.kjb
 completes successfully
etl/spec/dwh/pre_flight/pre_flight_ok_spec.kjb
  completes successfully
etl/spec/dwh/stage/stage_spec.kjb
  completes successfully
etl/spec/dwh/validate_params/validate_params_spec.kjb
  completes successfully
etl/spec/environment/mysql_driver_spec.kjb
  completes successfully
etl/spec/environment/robot_dir_is_set_spec.kjb - ROBOT_DIR variable is set correctly
  completes successfully
etl/spec/util/day_sequence/day_sequence_spec.kjb
  completes successfully
etl/spec/util/string_cleaner/string_cleaner_spec.kjb - identifiers are cleaned according to project standards
 completes successfully
```

etl/spec/daily/daily\_2001\_01\_01\_spec.kjb

etl/spec/daily/daily\_2017\_06\_20\_spec.kjb

etl/spec/daily/daily\_2017\_06\_20\_reload\_spec.kjb

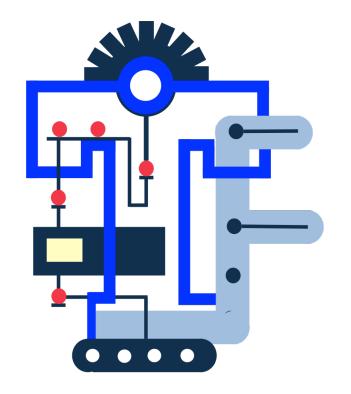
completes successfully

completes successfully

```
etl/spec/dummy/dummy/spec.kjb - verifies the test suite runs ETL as tests
   completes successfully
 etl/spec/dwh/load/load_spec.kjb
   completes successfully
 etl/spec/dwh/pre_flight/pre_flight_empty_file_spec.kjb
   completes successfully
 etl/spec/dwh/pre_flight/pre_flight_missing_file_spec.kjb
   completes successfully
 etl/spec/dwh/pre_flight/pre_flight_ok_spec.kjb
   completes successfully
 etl/spec/dwh/stage/stage_spec.kjb
   completes successfully
 etl/spec/dwh/validate_params/validate_params_spec.kjb
   completes successfully
 etl/spec/environment/mysql_driver_spec.kjb
   completes successfully
 etl/spec/environment/robot_dir_is_set_spec.kjb - ROBOT_DIR variable is set correctly
   completes successfully
 etl/spec/util/day_sequence/day_sequence_spec.kjb
   completes successfully
 etl/spec/util/string_cleaner/string_cleaner_spec.kjb - identifiers are cleaned according to project standards
   completes successfully
jdbc_helper
 dwh_db
   when selecting from steelwheels.customers
     has 126 customers
Finished in 1 minute 8.62 seconds (files took 1.25 seconds to load)
28 examples, 0 failures
```

COMPTERES SUCCESSIBLE





JRuby as a scripting language

# Jruby



Ruby language reference and std-lib documentation:

http://ruby-doc.org/

# Jruby



Jruby is a ruby language implementation on the JVM

http://jruby.org/

- excellent Java interop features
- developed and maintained by RedHat

## Jruby – A ruby primer



http://tryruby.org/

## Jruby — interactive ruby



```
$ bin/robot jruby -S jirb
irb(main):001:0> 1+2
=> 3
irb(main):002:0> "Hello World"
=> "Hello World"
irb(main):003:0>
```

## Jruby – variable assignment Twineworks



```
irb> greeting = "Hello"
=> "Hello"
irb> greeting
=> "Hello"
```

## Jruby – console output



```
irb> puts "Hello"
"Hello"
=> nil
```

## Jruby – calling methods



```
irb> "Hello".length
=> 5
irb> "Hello".length()
=> 5
irb> "Hello".gsub(/[aeiou]/, '*')
=> "H*]]*"
```

#### Jruby – conditionals



```
irb> if 5 > 3 then "Yes" else "No" end
=> "Yes"
irb> if 5 > 3
       "Yes"
     else
       "No"
     end
=> "Yes"
```

## Jruby – strings



```
irb> 'Hello World'
=> "Hello World"
irb> "Line1\nLine2"
=> "line1\nline2"
irb> puts "Line1\nLine2"
Line1
Line2
=> nil
```

## Jruby – strings



```
irb> answer = 42
=> 42
```

```
irb> "answer is: #{answer}"
```

=> "answer is: 42"

### Jruby – arrays



```
irb> []
=> []
irb> [12, 32, 42].size
=> 3
irb > items = [12, 32, 42]
=> [12, 32, 42]
irb> items.each do |x|
       puts x
     end
12
32
42
=> [12, 32, 42]
```

### Jruby – arrays



```
irb> items.each {|x| puts x}
12
32
42
=>[12, 32, 42]
```

### Jruby – arrays



```
irb> items.map \{|x| x+100\}
=> [112, 132, 142]
irb> [1, 2, 3, 4].reduce(0) {|a, x| a+x}
=> 10
irb> [1, 2, 3, 4].include? 0
=> false
irb> [1, 2, 3, 4].include? 3
=> true
```

### Jruby – hashes



```
irb> {}
=> {}
irb> {:first_name => "John", :last_name => "Doe"}
=> {:first_name => "John", :last_name => "Doe"}
irb> person = {:first_name => "John", :last_name => "Doe"}
=> {:first_name => "John", :last_name => "Doe"}
irb> person.each do |k, v|
       puts "#{k} is #{v}"
     end
first name is John
last name is Doe
=> {:first name=>"John", :last name=>"Doe"}
```

### Jruby – hashes



```
irb> person.keys
=> [:first_name, :last_name]
irb> person.values
=> ["John", "Doe"]
irb> person.to_a
=> [[:first_name, "John"], [:last_name, "Doe"]]
```

#### Jruby – environment vars



```
irb> ENV
=>{"ROBOT_DIR"=>"/Users/slawo/Desktop/etl-
testing/solution", "XPC_FLAGS"=>"0x0" ...}
irb> ENV.keys
=> ["ROBOT_DIR", "XPC_FLAGS",
"TERM_PROGRAM", "HOME", ...]
```

## Jruby – Ruby quick reference Twineworks

Ruby cheat sheets

https://github.com/marcusvmsa/cheatsheets/tree/master/ruby

# Jruby — running a ruby script Twineworks

\$ bin/robot jruby file

\$ bin/robot jruby ruby/hello.rb
Hello World!



Write a ruby script that outputs all environment variables to the console. Each variable goes on a sepearate line in the format: VAR=VALUE

```
$ bin/robot jruby ruby/environment_variables.rb
GEM_HOME=/Users/slawo/Desktop/etl-
testing/solution/bin/jruby/gem_home
TERM_PROGRAM_VERSION=388.1
USER=slawo
DISPLAY=/private/tmp/com.apple.launchd.TTWKUYpXJx/org.macosforge.xq
uartz:0
HOME=/Users/slawo
...
```

## Jruby – defining methods



```
irb> def add(a=0, b=0)
          a+b
     end
=> nil
irb> add
=> 0
irb> add 7
=> 7
irb> add 7, 9
=> 16
irb> add 7, 9, 6
ArgumentError: wrong number of arguments calling `add` (3 for 2)
```

## Jruby – command line args



```
irb> ARGV
=> []
irb> ARGV.size
=> 0
```

# Jruby – Dir



```
irb> Dir.pwd
=> "/Users/slawo/Desktop/etl-testing/solution"
irb> Dir.glob("./**/*.csv")
=> ["./bin/jruby/gem_home/gems/diff-lcs-1.3/spec/fixtures/ds1.csv",
 "./bin/jruby/gem_home/gems/diff-lcs-1.3/spec/fixtures/ds2.csv",
 "./data/in/2017/06/2017-06-18.csv",
 "./data/in/2017/06/2017-06-19.csv",
 "./data/in/2017/06/2017-06-20.csv",
 "./data/in/2017/06/2017-06-21.csv",
 "./etl/init/date/fixed_date_holidays.csv",
 "./etl/init/locale/iso_countries.csv",
 "./etl/spec/util/string_cleaner/string_cleaner_spec.csv",
```



Write a ruby script that outputs all file paths below a given directory that end in a given prefix.

The script prints a usage message if the number of arguments is not two.

```
$ bin/robot jruby ruby/find_in_dir.rb etl .kjb
etl/daily.kjb
etl/dwh/validate_params.kjb
etl/dwh/load/load.kjb
etl/dwh/pre_flight/pre_flight.kjb
...
$ bin/robot jruby ruby/find_in_dir.rb etl
usage: find_in_dir.rb dir suffix
```

## Jruby – metaprogramming



```
irb> class Hash
       def keys_reversed
         keys.reverse
       end
     end
=> nil
irb> {:a => "a", :b => "b"}.keys
=> [:a, :b]
irb> {:a => "a", :b => "b"}.keys_reversed
=> [:b, :a]
```

# Jruby – Gems



```
irb> "hello".blank?
NoMethodError: undefined method `blank?' for "hello":String
irb> require "active_support/all"
=> true
irb> "hello".blank?
=> false
irb> "".blank?
=> true
```

http://guides.rubyonrails.org/active support core extensions.html

## Jruby – Gems

# jruby/Gemfile



```
source 'http://rubygems.org'
gem 'gson'  # json processing
gem 'nokogiri' # xml processing
gem 'rspec' # testing framework
gem 'rspec-its'
gem 'rspec_junit_formatter' # rspec formatter for xml output on CI
gem 'jdbc-helper' # convenient jdbc queries
gem 'iniparse' # convenient ini parsing
gem 'activesupport', '4.2.8' # language core extensions
```



Add the cowsay gem <a href="https://github.com/johnnyt/cowsay">https://github.com/johnnyt/cowsay</a> to the project by modifying Gemfile and running

\$ bin/jruby/install-gems

Write a ruby script that has the avatar say the first argument, using the optional second argument for the template.

\$ bin/robot jruby ruby/cowsay.rb 'A good day to you!'



\$ bin/robot jruby ruby/cowsay.rb 'A good day to you!' bunny

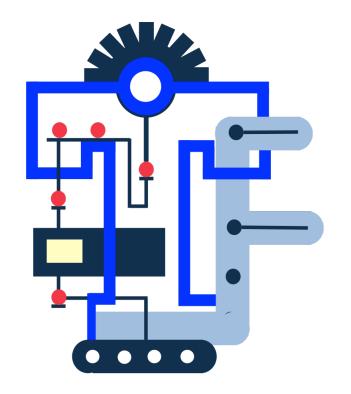
```
| A good day to you! |
      .( o ).
```



Write a ruby script that, given the path to a ktr or kjb file, outputs the description entered in the job or transformation properties dialog, if any.

- **\$ bin/robot jruby ruby/meta.rb etl/util/string\_cleaner.ktr** identifiers are cleaned according to project standards
- **\$ bin/robot jruby ruby/meta.rb etl/daily.kjb** main entry point for the robot application
- \$ bin/robot jruby ruby/meta.rb missing\_file
  file does not exist: missing\_file
- \$ bin/robot jruby ruby/meta.rb bin/load/help.txt unknown file format: bin/load/help.txt





Orchestrating the test suite with rspec

## Rspec



#### Rspec is a testing framework for ruby

http://rspec.info/

https://relishapp.com/rspec

### Rspec



#### \$ bin/robot test

- invokes rspec
- rspec loads spec/spec\_helper.rb
- spec/spec\_helper.rb sets configuration options
- includes all ruby files in spec/support
- ruby files in spec/support define any helper methods that the actual tests use
- rspec traverses the spec folder looking for files whose names end in \_spec.rb and loads them as tests

## Rspec - a spec file



```
describe "db clear command" do
  describe 'when db is not empty' do
    before :all do
      dwh db.load fixture 'spec/fixtures/steelwheels/steelwheels.sql'
     @out, @err, @status = execute "bin/robot db clear"
    end
    # print output of commands if anything failed
    after(:each) do |example|
     if example.exception
        puts @out
        puts ""
        puts @err
     end
    end
    it 'exits with exit code 0' do
      expect(@status).to eq 0
    end
    it "clears the db" do
      expect(dwh_db.query("SHOW TABLES").to_a.uniq.length).to eq 0
    end
  end
```

end

#### Rspec – describe blocks



Describe blocks are grouping tests into related units. They can be nested.

```
describe "some context" do
...
```

Reference documentation for the basic structure of a test file: <a href="https://relishapp.com/rspec/rspec-core/v/3-4/docs/example-groups/basic-structure-describe-it">https://relishapp.com/rspec/rspec-core/v/3-4/docs/example-groups/basic-structure-describe-it</a>

## Rspec – before and after



- Describe blocks can contain before and after hooks.
- If any tests within a describe block are run, the corresponding enclosing before and after hooks are executed as defined.
- The hooks can be scoped to run before/after all or each test.

```
describe "subsystem foo" do
  before :all do
    @my_var = "foo"
  end
  ...
end
```

## Rspec – before and after



In ETL testing the **before** hook is typically used to execute some ETL and capture any immediate output and exit codes into **instance** variables.

Instance variables are prepended with an **(a)** sign, and are available in before, after and example blocks.

https://www.relishapp.com/rspec/rspec-core/v/3-4/docs/hooks/before-and-after-hooks

#### Rspec – assertions



The actual assertions happen within it blocks, that describe expected behavior.

```
describe "subsystem foo" do

before :all do
   @my_var = "foo"
end

it "is equal to 'foo'" do
   expect(@my_var).to eq "foo"
end

end
```

Assertions are made through **expect**, which takes a value and supports a set of matchers. The matcher used above is **eq**.

### Rspec – assertions



Matcher reference documentation:

http://www.relishapp.com/rspec/rspec-expectations/v/3-4/docs/built-in-matchers

## Rspec – test execution



Rspec runs in two phases

Phase 1: collects tests, recording the structure as given by the describe blocks.

Phase 2: executes tests, using several command line options to limit the execution to just the tests fulfilling certain criteria, like having a certain substring in its name, or being tagged with certain keywords.

## Rspec – test execution



Run only tests containing the word 'clear' in their name or enclosing describe blocks:

\$ bin/robot test --example 'clear'

Run only tests tagged 'long\_running':

\$ bin/robot test --tag 'long\_running'

Run only tests in **spec/commands** 

\$ bin/robot test spec/commands

## Rspec – test execution



Reference documentation on how to tag a group of tests with metadata

<a href="https://relishapp.com/rspec/rspec-core/v/3-4/docs/metadata/user-defined-metadata">https://relishapp.com/rspec/rspec-core/v/3-4/docs/metadata/user-defined-metadata</a>

Reference documentation for rspec command line options <a href="https://relishapp.com/rspec/rspec-core/v/3-4/docs/command-line">https://relishapp.com/rspec/rspec-core/v/3-4/docs/command-line</a>

## Rspec helpers



#### spec/support/spec\_helpers.rb

```
def execute cmd
  stdin, stdout, stderr, thread = Open3.popen3(cmd)
  [stdout.read, stderr.read, thread.value.to_i]
end
```

Runs cmd as a shell command and returns an array returning [stdout, stderr, exit\_code]

## Rspec helpers



#### spec/support/spec\_helpers.rb

```
def dwh_db
    ...
end
```

Runs a jdbc database object offering convenient access to the **dwh** database as per <a href="https://github.com/junegunn/jdbc-helper">https://github.com/junegunn/jdbc-helper</a>

The database connection is created on demand, and closes automatically when rspec ends.

## Rspec helpers



#### spec/support/spec\_helpers.rb

```
def dwh_db
    ...
end
```

#### In addition

```
dwh_db.load_fixture(path) allows loading a sql or json fixture file
dwh_db.reset() triggers $ bin/robot db reset
```

## Running jobs as rspec tests



spec/etl/etl\_spec.rb

Recursively traverses **etl/spec** looking for files whose names end in **\_spec.kjb**, and dynamically generates a **describe** and **it** block for it.

Hence all such job files are part of the test suite.

## Test the load phase



Create a spec spec/load/load\_spec.rb

The spec verifies that /etl/dwh/load/load.kjb

- loads data from stage\_traffic
- has aggregated the records for same date, time, user and service
- has calculated the user\_tag by transforming it as expected
- has filled the service table as expected

Use /fixtures/load/2017-06-21.json to populate the stage\_traffic table.

## Test compliance



Create a spec
spec/environment/mysql\_driver\_spec.rb

The spec verifies that the file **\${ROBOT\_PDI\_HOME}/lib/mysql-connector-java-5.1.42-bin.jar** exists.

It also verifies that there is no other file matching pattern **mysql-connector-java-\*-bin.jar** in that folder.

#### **Declarative Tests**



Create a spec spec/daily/daily\_spec.rb

Recursively traverse **spec** looking for folders whose names end in **\_daily\_dec**, and dynamically generate **describe** and **it** blocks for them.

Each \_daily\_dec folder contains a folder input, a file run\_dates.txt, and a file expected.csv.

The spec starts by clearing the contents of the db.

Then, for each date in **run.dates**, the spec executes \$ bin/robot job run etl/daily.kjb -param:DATA\_DATE=<date> -param:DATA\_DIR=path/to/\*\_daily\_dec/input

Then the spec asserts that the state of the traffic tables is equal to expected.csv.

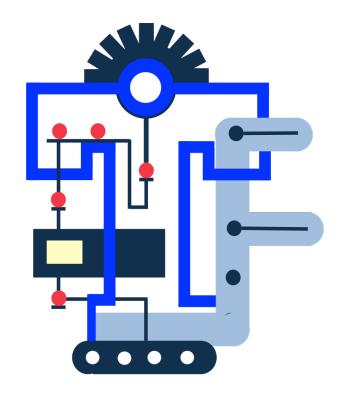
#### **Declarative Tests**



Convert all ETL functional tests from **etl/spec/daily** to the declarative format.

Place them in **spec/daily\_dec**.





Continuous testing with Jenkins

#### **Jenkins**





Jenkins is a continuous integration server.

It's basic role is to run the test suite and build any artifacts upon changes in version control.

Example server:

http://ci.pentaho.com/

## Installing Jenkins





Get **jenkins.war** from <a href="http://mirrors.jenkins.io/war/2.66/">http://mirrors.jenkins.io/war/2.66/</a>

\$ java -jar jenkins.war

Jenkins places all its configuration in ~/.jenkins

## Installing Jenkins





Install the following plugins:

AnsiColor
build-timeout-plugin
Credentials Binding Plugin
Email Extension Plugin
Git plugin
HTML Publisher plugin
Timestamper
Workspace Cleanup Plugin
Junit plugin

Some plugins will install as dependencies

# Installing Jenkins





Configure Jenkins to use a single executor.

# Placing the ETL in git



```
# create a bare repository
mkdir -p /home/slawo/Desktop/etl-project.git
cd /home/slawo/Desktop/etl-project.git
git init --bare --shared=group
# make the solution folder a git repository
cd /home/slawo/Desktop/solution
git init
git add .
git commit -m 'initial commit'
# potentially have to set up commit identity
git config --global user.email "slawomir.chodnicki@twineworks.com"
git config --global user.name "Slawomir Chodnicki"
# let local repository know where remote is
git remote add origin /home/slawo/Desktop/etl-project.git
# push to bare repository
git push -u origin master
```

### An environment for Jenkins Twineworks



Duplicate your **local** environment for jenkins as **test** 

Use {WORKSPACE} as placeholder for project directory in

- environments/test/.kettle/kettle.properties
- environments/test/environment

## Database configuration



```
$ mysql -u root

CREATE DATABASE dwh_jenkins;

GRANT ALL PRIVILEGES ON dwh_jenkins.* TO etl@'%' identified by 'password';

GRANT ALL PRIVILEGES ON dwh_jenkins.* TO etl@'localhost' identified by 'password';

FLUSH PRIVILEGES;
```

## Main test job



```
# bail on errors
set -e
# set environment to use
export ROBOT_ENVIRONMENT=test
# set workspace directory in config files
sed -i.bak s:{WORKSPACE}:${WORKSPACE}:g environments/test/.kettle/kettle.properties
sed -i.bak s:{WORKSPACE}:${WORKSPACE}:g environments/test/environment
# install dependencies
bin/jruby/install
# run test suite
bin/robot test
```

### Ensure Jenkins can run tests



Get the test suite to pass on Jenkins, fixing configuration issues as they arise.

Don't forget to properly **commit** and **push** your changes.

## Jenkins artefacts





Configure Jenkins to use the **rspec-integration.xml** test report file.

## Jenkins artefacts





Configure Jenkins to **archive** all artefacts on successful builds.

# Jenkins nightly





Configure Jenkins to create a nightly build in a separate job.

### Jenkins concurrent builds





What would it take to enable multiple executors?

# Jenkins pipelines



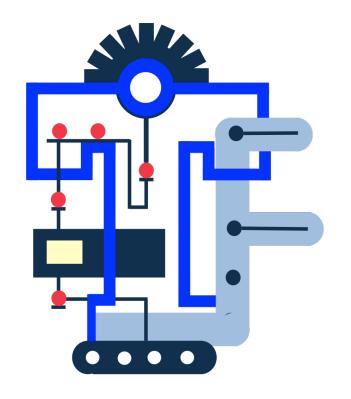


Jenkins pipelines implement workflows.

The Blue Ocean project provides a sensible UI for pipelines.

https://jenkins.io/projects/blueocean/





Simple deployment from git

# Deploying from git



Tag two releases as

```
$ git tag releases/2017-06-21_r1 commit_hash
```

\$ git tag releases/2017-06-21\_r2 commit\_hash

And push to shared repo

```
$ git push origin releases/2017-06-21_r1
```

\$ git push origin releases/2017-06-21\_r2

List remote tags

\$git Is-remote --tags origin

## Example deployment scripts Twineworks

```
-- conf
     # configuration
-- activate.sh
          # symlinks a tag as active
```

## Example deployment scripts Twineworks



```
Extract and activate 2017-06-21 r1
```

- \$./extract.sh 2017-06-21 r1
- \$ ./activate.sh 2017-06-21\_r1

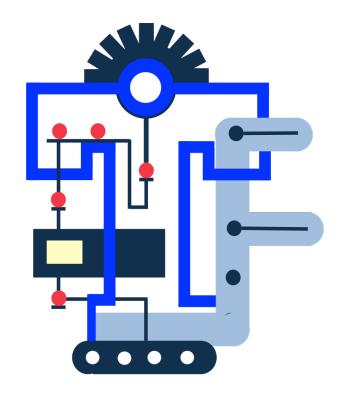
#### Extract and activate 2017-06-21 r2

- \$ ./extract.sh 2017-06-21\_r2
- \$./activate.sh 2017-06-21 r2

#### Rollback to **2017-06-21 r1**

\$./activate.sh 2017-06-21 r1





Database schema migration

## Manual migration scripts



Manually perform DB schema migrations by creating a SQL script or kjb to migrate to the next version.

After deployment run

\$ bin/robot db run etl/migrations/release\_panda/add\_tables.sql

Or if it is a kjb

\$ bin/robot job run etl/migrations/release\_panda/panda\_migration.kjb

Test and dry run migrations on QA like any other code.

Avoid destructive changes, i.e. rename columns instead of dropping them. Drop after confirmed lasting release success.

# Managed migration scripts

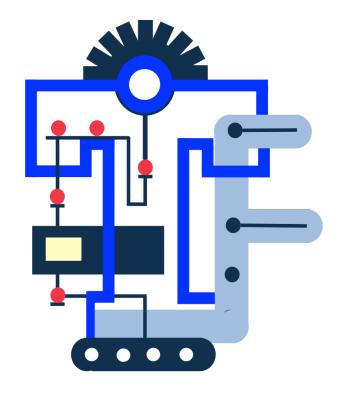


Flyway automates DB schema migrations by storing a version number in the schema and maintaining a set of SQL scripts or java code to migrate from one version to the next.



https://flywaydb.org/





Thank you!